

## EXP\_21

The screenshot shows a Jupyter Notebook interface with two code cells and their corresponding outputs.

**Code Cell 1:**

```
import numpy as np
import pandas as pd
from scipy import stats

# Load dataset safely
df = pd.read_csv(
    "mining_projects.csv",
    engine="python",
    on_bad_lines="skip"
)

# Convert all columns to numeric where possible
numeric_df = df.apply(pd.to_numeric, errors='coerce')

# Drop columns that are fully NaN
numeric_df = numeric_df.dropna(axis=1, how='all')

# Convert to NumPy array and flatten
concentration_data = numeric_df.values.flatten()

# Remove NaN values
concentration_data = concentration_data[~np.isnan(concentration_data)]

# Check data availability
if len(concentration_data) == 0:
    raise ValueError("No numeric concentration data found in dataset.")

# User inputs
sample_size = int(input("Enter sample size: "))
confidence_level = float(input("Enter confidence level (e.g., 0.95): "))
precision = float(input("Enter desired precision (margin of error): "))

# Random sample
sample = np.random.choice(concentration_data, size=sample_size, replace=False)

# Point estimation
mean_estimate = np.mean(sample)
std_dev = np.std(sample, ddof=1)

# Z-score
z = stats.norm.ppf((1 + confidence_level) / 2)

# Margin of error
margin_error = z * (std_dev / np.sqrt(sample_size))

# Confidence interval
lower = mean_estimate - margin_error
upper = mean_estimate + margin_error

# Output
print("\n--- Estimation Results ---")
print("Point Estimate (Mean): {:.4f}".format(mean_estimate))
print("({:.4f}, {:.4f})% Confidence Interval: ({:.4f}, {:.4f})".format(int(confidence_level*100), lower, upper))
print("Margin of Error: {:.4f}".format(margin_error))

if margin_error <= precision:
    print("Desired precision achieved.")
else:
    print("Desired precision NOT achieved.")
```

**Code Cell 2:**

```
if margin_error <= precision:
    print("Desired precision achieved.")
else:
    print("Desired precision NOT achieved.")
```

**Output:**

```
*** Enter sample size: 5
Enter confidence level (e.g., 0.95): 0.95
Enter desired precision (margin of error): 56

--- Estimation Results ---
Point Estimate (Mean): 1232.6000
95% Confidence Interval: (287.9373, 2177.2627)
Margin of Error: 944.6627
Desired precision NOT achieved.
```

## EXP\_22

**EXP\_22** Imagine you are an analyst for a popular online shopping website. Your task is to analyze customer reviews and provide insights on the average rating.

```

import pandas as pd
import numpy as np
from scipy import stats

# Load dataset safely (important fix)
df = pd.read_csv(
    "customer_reviews.csv",
    engine="python",
    on_bad_lines="skip"
)

# Extract and clean rating column
ratings = pd.to_numeric(df["reviews.rating"], errors="coerce").dropna()

# Sample statistics
sample_mean = ratings.mean()
sample_std = ratings.std(ddof=1)
sample_size = len(ratings)

# User input
confidence_level = float(input("Enter confidence level (e.g., 0.95):"))

# Z-score
z = stats.norm.ppf((1 + confidence_level) / 2)

# Margin of error
margin_error = z * (sample_std / np.sqrt(sample_size))

```

**EXP\_22** Imagine you are an analyst for a popular online shopping website. Your task is to analyze customer reviews and provide insights on the average rating.

```

# Margin of error
margin_error = z * (sample_std / np.sqrt(sample_size))

# Confidence interval
lower = sample_mean - margin_error
upper = sample_mean + margin_error

# Output
print("\n--- Customer Review Analysis ---")
print(f"Average Rating (Sample Mean): {sample_mean:.2f}")
print(f"95% Confidence Interval: ({lower:.2f}, {upper:.2f})")

# Customer satisfaction level
if sample_mean >= 4:
    print("Customer Satisfaction Level: High")
elif sample_mean >= 3:
    print("Customer Satisfaction Level: Moderate")
else:
    print("Customer Satisfaction Level: Low")

```

\*\*\* Enter confidence level (e.g., 0.95): 0.95

--- Customer Review Analysis ---  
Average Rating (Sample Mean): 4.60  
95% Confidence Interval: (4.58, 4.62)  
Customer Satisfaction Level: High

## EXP\_23

**EXP\_23** To Analyze the data using hypothesis testing and calculate the p-value to determine if the new treatment has a statistically significant effect compared to the control group.

```

import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import ttest_ind

df = pd.read_csv("diabetes.csv")

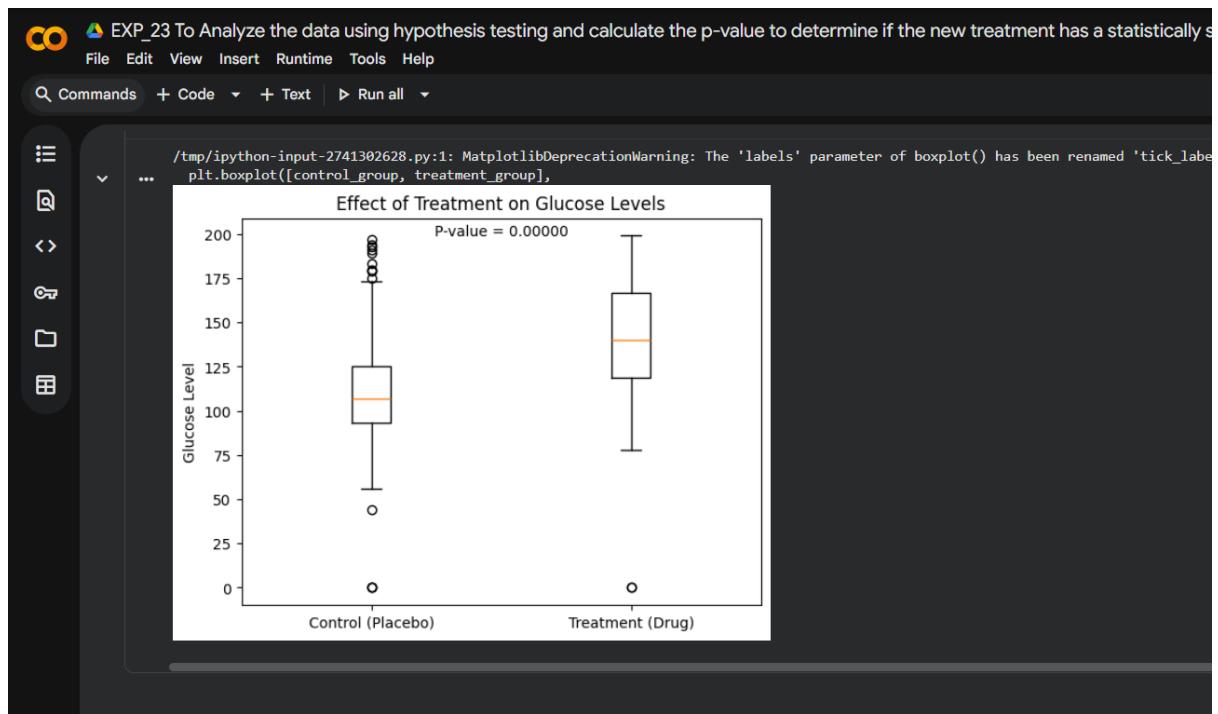
control_group = df[df["Outcome"] == 0]["Glucose"]
treatment_group = df[df["Outcome"] == 1]["Glucose"]

t_stat, p_value = ttest_ind(control_group, treatment_group)
print("T-statistic:", t_stat)
print("P-value:", p_value)

T-statistic: -14.600000000059738
P-value: 8.935431645289912e-43

plt.boxplot([control_group, treatment_group],
            labels=["Control (Placebo)", "Treatment (Drug)"])
plt.title("Effect of Treatment on Glucose Levels")
plt.ylabel("Glucose Level")
plt.text(1.5, max(df["Glucose"]), f"P-value = {p_value:.5f}", ha="center")
plt.show()

```



## EXP\_24

EXP\_24 Write a Python program that allows the user to input the features of a new patient and the value of k (number of neighbors).

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler

df = pd.read_csv("diabetes.csv")

X = df.drop("Outcome", axis=1)
y = df["Outcome"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
k = int(input("Enter value of k (neighbors): "))

print("\nEnter patient details:")
patient = []
for col in X.columns:
    val = float(input(f"{col}: "))
    patient.append(val)

```

EXP\_24 Write a Python program that allows the user to input the features of a new patient and the value of k (number of neighbors). ⭐ ⚙

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▶

```
[ ] patient.append(val)

[ ] Enter value of k (neighbors): 5

[ ] Enter patient details:
Pregnancies: 2
Glucose: 8
BloodPressure: 5
SkinThickness: 5
Insulin: 3
BMI: 8
DiabetesPedigreeFunction: 5
Age: 45

[ ] knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)
patient_df = pd.DataFrame([patient], columns=X.columns)
patient_scaled = scaler.transform(patient_df)
prediction = knn.predict(patient_scaled)

[ ] if prediction[0] == 1:
    print("\nPrediction: Patient HAS the medical condition")
else:
    print("\nPrediction: Patient DOES NOT have the medical condition")

[ ] Prediction: Patient DOES NOT have the medical condition
```

## EXP\_25

EXP\_25 Write a Python program that loads the Iris dataset from scikit-learn, and allows the user to input the sepal length, sepal width, petal length, and pe...

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▶

```
[ ] import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score

[ ] columns = [
    "sepal_length",
    "sepal_width",
    "petal_length",
    "petal_width",
    "species"
]

[ ] df = pd.read_csv("iris.csv", header=None, names=columns)

[ ] X = df.iloc[:, 0:4]
y = df["species"]
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

EXP\_25 Write a Python program that loads the Iris dataset from scikit-learn, and allows the user to input the sepal length, sepal width, petal length, and petal width of a new flower, and predicts its species.

```
model = DecisionTreeClassifier()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Model Accuracy:", accuracy)

Model Accuracy: 1.0

print("\nEnter new flower measurements:")
sl = float(input("Sepal Length: "))
sw = float(input("Sepal Width: "))
pl = float(input("Petal Length: "))
pw = float(input("Petal Width: "))

new_flower = pd.DataFrame(
    [[sl, sw, pl, pw]],
    columns=["sepal_length", "sepal_width", "petal_length", "petal_width"]
)

prediction = model.predict(new_flower)

print("\nPredicted Iris Species:", prediction[0])
```

EXP\_25 Write a Python program that loads the Iris dataset from scikit-learn, and allows the user to input the sepal length, sepal width, petal length, and petal width of a new flower, and predicts its species.

```
)  
prediction = model.predict(new_flower)  
print("\nPredicted Iris Species:", prediction[0])  
  
Enter new flower measurements:  
Sepal Length: 5.4  
Sepal Width: 3.5  
Petal Length: 5.6  
Petal Width: 2.6  
  
Predicted Iris Species: Iris-virginica
```

## EXP\_26

EXP\_26 Write a Python program that allows the user to input the features (area, number of bedrooms, etc.) of a new house.

```
import pandas as pd
from sklearn.linear_model import LinearRegression

df = pd.read_csv("housing.csv")
print(df.columns)

Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
       'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
       'parking', 'prefarea', 'furnishingstatus'],
      dtype='object')

X = df[['area', 'bedrooms', 'bathrooms']]
y = df['price']

model = LinearRegression()
model.fit(X, y)
print("Enter details of the new house:")
area = float(input("Area (in sq.ft): "))
bedrooms = int(input("Number of bedrooms: "))
bathrooms = int(input("Number of bathrooms: "))

Enter details of the new house:
Area (in sq.ft): 3000
Number of bedrooms: 4
Number of bathrooms: 3
```

EXP\_26 Write a Python program that allows the user to input the features (area, number of bedrooms, etc.) of a new house. ⭐ ☁

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

```
[ ] bedrooms = int(input("Number of bedrooms: "))
bathrooms = int(input("Number of bathrooms: "))

[+] Enter details of the new house:
Area (in sq.ft): 3000
Number of bedrooms: 4
Number of bathrooms: 3

[ ] new_house = pd.DataFrame(
    [[area, bedrooms, bathrooms]],
    columns=['area', 'bedrooms', 'bathrooms']
)

predicted_price = model.predict(new_house)

print("\nPredicted House Price:", predicted_price[0])

[+] Predicted House Price: 6748545.281572397
```

## EXP\_27

EXP\_27 Write a Python program that allows the user to input the features (e.g., usage minutes, contract duration) of a new customer. ⭐ ☁

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

```
[+] import pandas as pd
from sklearn.linear_model import LogisticRegression

[+] df = pd.read_csv("telco_churn.csv")

[+] print(df['Churn'].value_counts())
Churn
No      5174
Yes     1869
Name: count, dtype: int64

[+] X = df[['tenure', 'MonthlyCharges', 'TotalCharges']].copy()
y = df['Churn']

[+] X['TotalCharges'] = pd.to_numeric(X['TotalCharges'], errors='coerce')

X = X.fillna(X.mean(numeric_only=True))
```

EXP\_27 Write a Python program that allows the user to input the features (e.g., usage minutes, contract duration) of a new customer. ⭐ ⌂

```
File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all

model = LogisticRegression(max_iter=1000)
model.fit(X, y)

# User input
print("\nEnter new customer details:")
tenure = float(input("Tenure (months): "))
monthly_charges = float(input("Monthly Charges: "))
total_charges = float(input("Total Charges: "))

# Predict churn
new_customer = pd.DataFrame(
    [[tenure, monthly_charges, total_charges]],
    columns=['tenure', 'MonthlyCharges', 'TotalCharges']
)

Enter new customer details:
Tenure (months): 6
Monthly Charges: 250
Total Charges: 8000

prediction = model.predict(new_customer)

if prediction[0] == 1:
    print("\nPrediction: Customer WILL churn")
else:
    print("\nPrediction: Customer will NOT churn")

Prediction: Customer will NOT churn
```

## EXP\_28

EXP\_28 Write a Python program that allows the user to input the shopping-related features of a new customer. ⭐ ⌂

```
File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all

import pandas as pd
from sklearn.cluster import KMeans

df=pd.read_csv("Mall_Customers.csv")

X = df[['Annual Income (k$)', 'Spending Score (1-100)']]

# Train K-Means model
kmeans = KMeans(n_clusters=5, random_state=42)
kmeans.fit(X)

KMeans(n_clusters=5, random_state=42)
```

EXP\_28 Write a Python program that allows the user to input the shopping-related features of a new customer. ⭐ ↗

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

```
[ ] print("Enter new customer details:")
income = float(input("Annual Income (in k$): "))
spending_score = float(input("Spending Score (1-100): "))

new_customer = pd.DataFrame(
    [[income, spending_score]],
    columns=['Annual Income (k$)', 'Spending Score (1-100)']
)

cluster = kmeans.predict(new_customer)

print("\nThe new customer belongs to Cluster:", cluster[0])
```

Enter new customer details:  
Annual Income (in k\$): 60  
Spending Score (1-100): 65

The new customer belongs to Cluster: 0

## EXP\_29

EXP\_29 Write a Python program that loads a dataset and trained model from scikit-learn and it should calculate the F1-Score,Recall,Precision,Accuracy.

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

```
[ ] import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
[ ] df = pd.read_csv("dataset.csv")
print(df.head())
print(df.columns)
```

	Id	Cl.thickness	Cell.size	Cell.shape	Marg.adhesion	Epith.c.size	\
0	1000025	5	1	1	1	2	
1	1002945	5	4	4	5	7	
2	1015425	3	1	1	1	2	
3	1016277	6	8	8	1	3	
4	1017023	4	1	1	3	2	

```
[ ] Bare.nuclei Bl.cromatin Normal.nucleoli Mitoses Class
0 1.0 3 1 1 0
1 10.0 3 2 1 0
2 2.0 3 1 1 0
3 4.0 3 7 1 0
4 1.0 3 1 1 0
Index(['Id', 'Cl.thickness', 'Cell.size', 'Cell.shape', 'Marg.adhesion',
       'Epith.c.size', 'Bare.nuclei', 'Bl.cromatin', 'Normal.nucleoli',
       'Mitoses', 'Class'],
      dtype='object')
```

**EXP\_29** Write a Python program that loads a dataset and trained model from scikit-learn and it should calculate the F1-Score,Recall,Precision,Accuracy.

```

File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all

[ ] features = input("\nEnter feature column names (comma separated): ").split(",")
features = [f.strip() for f in features]

target = input("Enter target column name: ").strip()

# Prepare data
X = df[features]
y = df[target]

Enter feature column names (comma separated): Cl.thickness,Cell.size,Cell.shape,Marg.adhesion,Epith.c.size,Bare.nuclei,Bl.cromatin,Normal.nucleoli,Mitoses
Enter target column name: Class

[ ] # Handle missing values
X = X.fillna(X.mean())

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Train model
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

```

```

# Evaluation metrics
print("\nModel Evaluation Metrics:")
print("Accuracy : ", accuracy_score(y_test, y_pred))
print("Precision: ", precision_score(y_test, y_pred))
print("Recall   : ", recall_score(y_test, y_pred))
print("F1 Score : ", f1_score(y_test, y_pred))

Model Evaluation Metrics:
Accuracy : 0.9571428571428572
Precision: 0.975609756097561
Recall   : 0.8888888888888888
F1 Score : 0.9302325581395349

```

## EXP\_30

**EXP\_30** Write a Python program that loads the car dataset and allows the user to input the features of a new car they want to sell.

```

File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all

[ ] import pandas as pd
from sklearn.tree import DecisionTreeRegressor
from sklearn.preprocessing import LabelEncoder

[ ] df = pd.read_csv("car_price.csv")
print(df.columns)

Index(['Car_Name', 'Year', 'Selling_Price', 'Present_Price', 'Kms_Driven',
       'Fuel_Type', 'Seller_Type', 'Transmission', 'Owner'],
      dtype='object')

[ ] df = df[['Selling_Price', 'Year', 'Kms_Driven', 'Fuel_Type', 'Transmission']]

# Encode categorical features
le_fuel = LabelEncoder()

```

EXP\_30 Write a Python program that loads the car dataset and allows the user to input the features of a new car they want to sell.

```
File Edit View Insert Runtime Tools Help
Commands + Code + Text ▶ Run all ▾

df['Fuel_Type'] = le_fuel.fit_transform(df['Fuel_Type'])
df['Transmission'] = le_trans.fit_transform(df['Transmission'])

# Features and target
X = df[['Year', 'Kms_Driven', 'Fuel_Type', 'Transmission']]
y = df['Selling_Price']

# Train CART model (Decision Tree Regression)
model = DecisionTreeRegressor(random_state=42)
model.fit(X, y)

DecisionTreeRegressor(random_state=42)

print("\nEnter details of the new car:")
year = int(input("Manufacturing year: "))
kms = float(input("Kilometers driven: "))
fuel = input("Fuel type (Petrol/Diesel/CNG): ")
trans = input("Transmission (Manual/Automatic): ")

# Encode user inputs
fuel_enc = le_fuel.transform([fuel])[0]
trans_enc = le_trans.transform([trans])[0]
```

EXP\_30 Write a Python program that loads the car dataset and allows the user to input the features of a new car they want to sell.

```
File Edit View Insert Runtime Tools Help
Commands + Code + Text ▶ Run all ▾

# Encode user inputs
fuel_enc = le_fuel.transform([fuel])[0]
trans_enc = le_trans.transform([trans])[0]

# Predict price (clean way)
new_car = pd.DataFrame(
    [[year, kms, fuel_enc, trans_enc]],
    columns=['Year', 'Kms_Driven', 'Fuel_Type', 'Transmission']
)

predicted_price = model.predict(new_car)
print("\nPredicted Car Price:", predicted_price[0])

Enter details of the new car:
Manufacturing year: 2018
Kilometers driven: 40000
Fuel type (Petrol/Diesel/CNG): Petrol
Transmission (Manual/Automatic): Manual

Predicted Car Price: 3.5

node_indicator = model.decision_path(new_car)
feature = model.tree_.feature
threshold = model.tree_.threshold

print("\nDecision Path (Conditions followed):")
for node_id in node_indicator.indices:
    if feature[node_id] != -2:
        print(f"\t{X.columns[feature[node_id]]} <= {threshold[node_id]}")
```

```
...
Decision Path (Conditions followed):
Fuel_Type <= 1.5
Year <= 2014.5
Kms_Driven <= 6800.0
Transmission <= 0.5
Kms_Driven <= 10990.0
Kms_Driven <= 11824.5
Kms_Driven <= 12239.5
Kms_Driven <= 15000.5
Kms_Driven <= 16101.0
Kms_Driven <= 18398.0
Kms_Driven <= 49781.0
Kms_Driven <= 49164.5
Kms_Driven <= 22835.5
Kms_Driven <= 24339.0
Kms_Driven <= 24900.0
Year <= 2015.5
Kms_Driven <= 30111.5
Kms_Driven <= 36744.0
```