

DAY 6 LAB EXPERIMENTS

Name: L Aravind

Reg No : 192424080

EXP _ 31

Welcome to Colab Cannot save changes

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Copy to Drive

Files

- ..
- sample_data
- customer_segmentation_data_5..

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

df = pd.read_csv("customer_segmentation_data_500.csv")

features = df[['Annual_Spend', 'Visits', 'Time_Spent', 'Age']]

scaler = StandardScaler()
scaled_features = scaler.fit_transform(features)

kmeans = KMeans(n_clusters=4, random_state=42)
df['Cluster'] = kmeans.fit_predict(scaled_features)

print(df.head(10))

plt.figure()
plt.scatter(
    df['Annual_Spend'],
    df['Visits'],
    c=df['Cluster']
)
plt.xlabel("Annual Spend")
plt.ylabel("Number of Visits")
plt.title("Customer Segmentation using K-Means")
plt.show()

# Step 8: Analyze cluster characteristics
cluster_summary = df.groupby('Cluster').mean()
print("InCluster Summary:")
print(cluster_summary)
```

	Customer_ID	Annual_Spend	Visits	Time_Spent	Age	Cluster
0	1	25795	39	129	46	3
1	2	10860	55	115	49	0
2	3	48158	14	63	41	2
3	4	54732	21	37	31	2
4	5	21284	25	69	31	2
5	6	16265	35	112	55	0
6	7	26850	35	71	54	0
7	8	47194	57	70	44	0
8	9	31962	30	130	42	3
9	10	57191	48	35	25	1

Customer Segmentation using K-Means

Cluster Summary:

Cluster	Customer_ID	Annual_Spend	Visits	Time_Spent	Age
0	226.050420	36054.109244	46.537815	85.159664	48.630252
1	267.174242	33028.348485	47.196970	73.106061	28.954545
2	253.671875	33857.312500	23.718750	51.609375	40.062500
3	253.000000	33107.793388	22.702479	108.173554	37.280992

EXP_32

The screenshot shows a Jupyter Notebook interface with a sidebar on the left displaying a file tree. The tree includes a folder named 'sample_data' containing 'customer_segmentation_data_5..csv' and 'house_price_dataset_500.csv'. The main area contains Python code for performing a linear regression analysis on house price data.

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

df = pd.read_csv("house_price_dataset_500.csv")

print("Dataset Preview:")
print(df.head(5))

plt.figure()
plt.scatter(df['House_Size_sqft'], df['House_Price'])
plt.xlabel("House Size (sq.ft)")
plt.ylabel("House Price")
plt.title("Bivariate Analysis: House Size vs House Price")
plt.show()

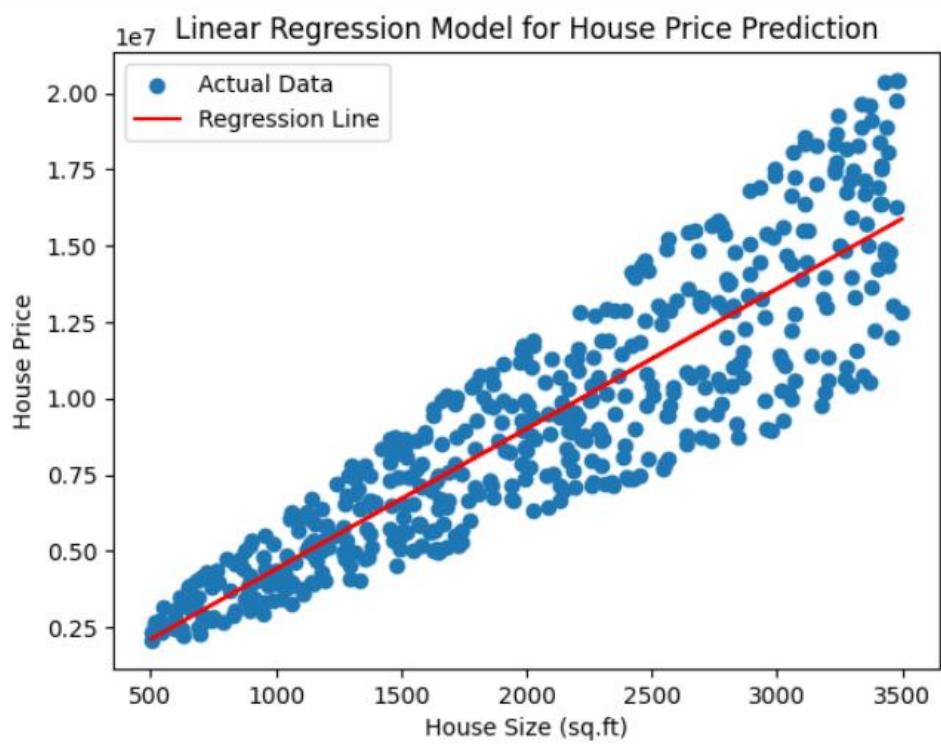
X = df[['House_Size_sqft']] # Independent variable
y = df['House_Price'] # Dependent variable

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("\nModel Evaluation:")
print("Mean Squared Error (MSE):", mse)
print("R^2 Score:", r2)
plt.figure()
plt.scatter(X, y, label="Actual Data")
plt.plot(X, model.predict(X), color='red', label="Regression Line")
plt.xlabel("House Size (sq.ft)")
plt.ylabel("House Price")
plt.title("Linear Regression Model for House Price Prediction")
plt.legend()
plt.show()
```

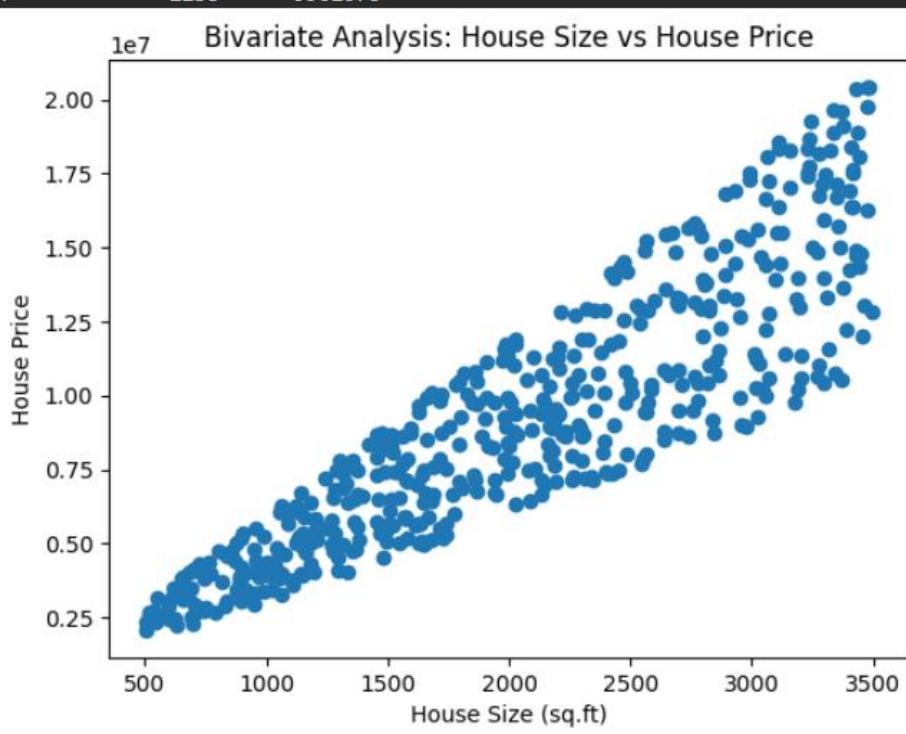
Disk 86.55 GB available

Model Evaluation:
Mean Squared Error (MSE): 3815179089697.245
 R^2 Score: 0.7897057805416927



Dataset Preview:

```
...> House_Size_sqft  House_Price
0           1360      6502160
1           1794     10493106
2           1630      9468670
3           1595      8754955
4           2138      9901078
```



EXP_33

The screenshot shows a Jupyter Notebook interface with two code cells and their corresponding outputs.

Code Cell 1:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
df = pd.read_csv("car_price_dataset_500.csv")

print(df.head())
X = df[['Engine_Size', 'Horsepower', 'Mileage']]
y = df['Price']
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print("Mean Squared Error (MSE):", mean_squared_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))
importance = pd.DataFrame({
    "Feature": X.columns,
    "Coefficient": model.coef_
})

print("\nFeature Importance:")
print(importance)
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual vs Predicted Car Prices")
plt.show()
```

Code Cell 2 (Output):

```
...   Engine_Size Horsepower Mileage  Price
0      3974        185     10  1635358
1      1660        177     24  646392
2      2094        107     11  662757
3      1930        148     23  629294
4      1895        296     17  641504
Mean Squared Error (MSE): 27869509320.7804
R2 Score: 0.7444798415577418

Feature Importance:
   Feature  Coefficient
0 Engine_Size  301.540411
1 Horsepower   803.342570
2 Mileage     -1716.511132
```

Figure:

A scatter plot titled "Actual vs Predicted Car Prices". The x-axis is labeled "Actual Price" and ranges from 0.4 to 1.8 million. The y-axis is labeled "Predicted Price" and ranges from 0.4 to 1.4 million. The plot shows a dense cloud of blue data points forming a strong positive linear trend, indicating a good fit of the linear regression model.

EXP_34

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report
df = pd.read_csv("medical_knn_dataset_500.csv")

print(df.head())
X = df[['Age', 'Gender', 'Blood_Pressure', 'Cholesterol']]
y = df['Outcome']
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print("\nAccuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred, pos_label="Good"))
print("Recall:", recall_score(y_test, y_pred, pos_label="Good"))
print("F1 Score:", f1_score(y_test, y_pred, pos_label="Good"))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
results = pd.DataFrame({
    "Actual": y_test.values,
    "Predicted": y_pred
})

print("\nPrediction Results:")
print(results.head(10))
```

The screenshot shows a Jupyter Notebook interface with two panes. The left pane displays a file tree with several CSV files: sample_data, car_price_dataset_500.csv, customer_segmentation_data_500.csv, house_price_dataset_500.csv, and medical_knn_dataset_500.csv. The right pane shows the execution results of the provided Python code.

Execution results:

- Printed dataset head:

	Age	Gender	Blood_Pressure	Cholesterol	Outcome
0	58	1	166	259	Bad
1	71	0	122	273	Bad
2	48	0	90	236	Bad
3	34	1	110	204	Good
4	62	0	144	181	Bad

- Printed classification metrics:

	accuracy	Precision	Recall	F1 Score
accuracy	0.97			
Precision	0.90625			
Recall	1.0			
F1 Score	0.9508196721311475			

- Printed Classification Report:

	precision	recall	f1-score	support
Bad	1.00	0.96	0.98	71
Good	0.91	1.00	0.95	29
accuracy			0.97	100
macro avg	0.95	0.98	0.96	100
weighted avg	0.97	0.97	0.97	100

- Printed Prediction Results:

	Actual	Predicted
0	Bad	Bad
1	Bad	Bad
2	Bad	Bad
3	Good	Good
4	Good	Good
5	Good	Good
6	Bad	Bad
7	Bad	Bad
8	Good	Good
9	Good	Good

EXP_35

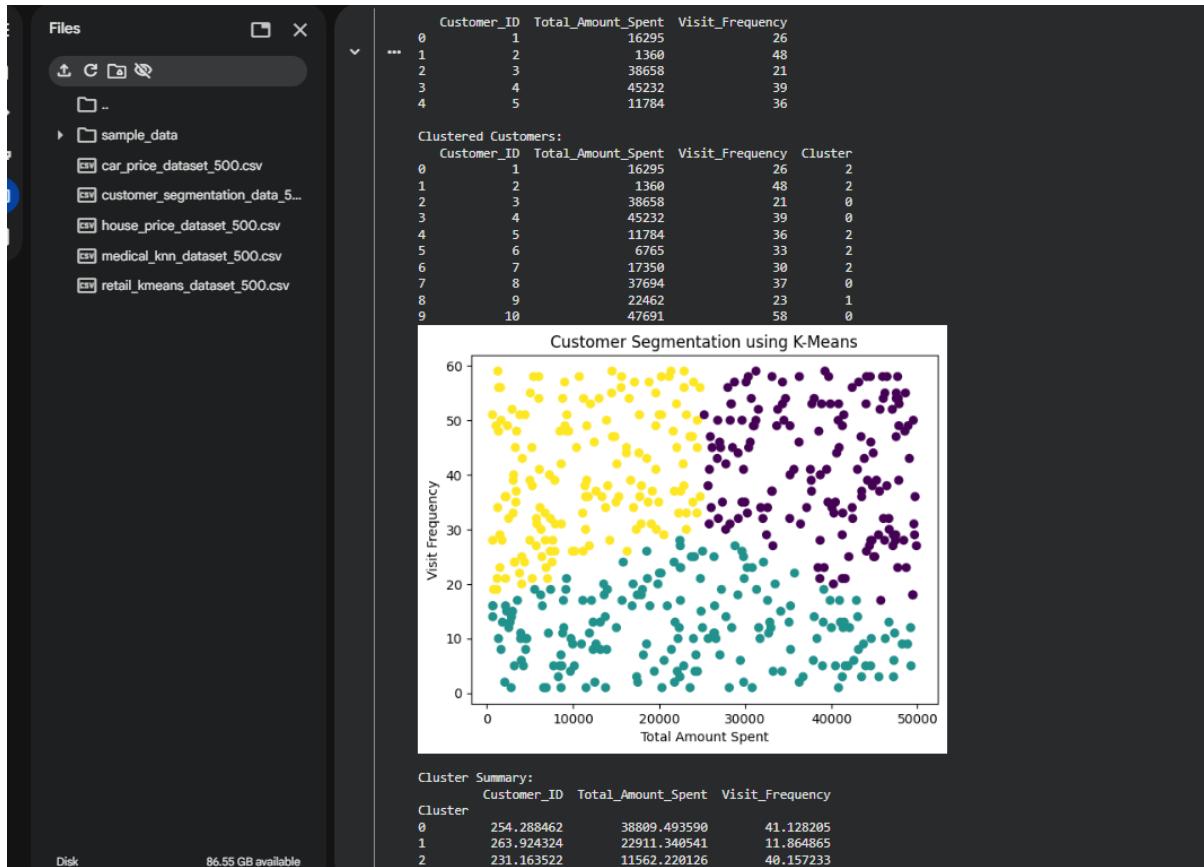
```
7   Bad    Bad
8   Good   Good
9   Good   Good

[5]  Os
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

df = pd.read_csv("retail_kmeans_dataset_500.csv")

print(df.head())
X = df[['Total_Amount_Spent', 'Visit_Frequency']]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

kmeans = KMeans(n_clusters=3, random_state=42)
df['Cluster'] = kmeans.fit_predict(X_scaled)
print("\nClustered Customers:")
print(df.head(10))
plt.scatter(
    df['Total_Amount_Spent'],
    df['Visit_Frequency'],
    c=df['Cluster']
)
plt.xlabel("Total Amount Spent")
plt.ylabel("Visit Frequency")
plt.title("Customer Segmentation using K-Means")
plt.show()
cluster_summary = df.groupby('Cluster').mean()
print("\nCluster Summary:")
print(cluster_summary)
```



EXP_36

```
1      263.924324    22911.340541    11.864865
2      231.163522    11562.220126    40.157233
```

[7]

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

df = pd.read_csv("stock_price_variability_500.csv")

print(df.head())

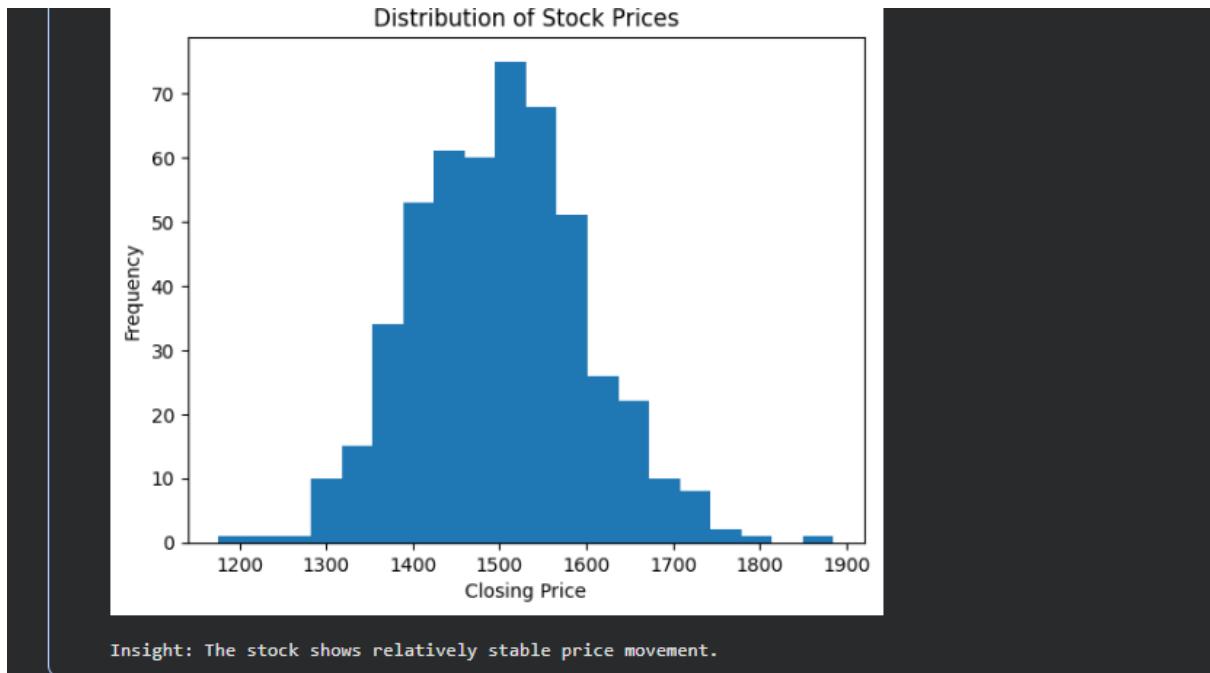
mean_price = df['Closing_Price'].mean()
variance_price = df['Closing_Price'].var()
std_dev_price = df['Closing_Price'].std()

print("\nStock Price Statistics:")
print("Mean Price:", mean_price)
print("Variance:", variance_price)
print("Standard Deviation:", std_dev_price)
plt.figure()
plt.plot(df['Day'], df['Closing_Price'])
plt.xlabel("Trading Day")
plt.ylabel("Closing Price")
plt.title("Stock Price Movement Over Time")
plt.show()
plt.figure()
plt.hist(df['Closing_Price'], bins=20)
plt.xlabel("Closing Price")
plt.ylabel("Frequency")
plt.title("Distribution of Stock Prices")
plt.show()
# -----
if std_dev_price > 100:
    print("\nInsight: The stock shows high price volatility.")
else:
    print("\nInsight: The stock shows relatively stable price movement.")
```

...	Day	Closing_Price
0	1	1549.67
1	2	1486.17
2	3	1564.77
3	4	1652.30
4	5	1476.58

Stock Price Statistics:
Mean Price: 1500.6839600000003
Variance: 9628.564469657722
Standard Deviation: 98.12524888966

Stock Price Movement Over Time



EXP_37

```

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
df = pd.read_csv("study_time_exam_score_500.csv")

print(df.head())
correlation = df['Study_Time_Hours'].corr(df['Exam_Score'])
print("\nCorrelation between Study Time and Exam Score:", correlation)

plt.figure()
plt.scatter(df['Study_Time_Hours'], df['Exam_Score'])
plt.xlabel("Study Time (Hours)")
plt.ylabel("Exam Score")
plt.title("Scatter Plot: Study Time vs Exam Score")
plt.show()

plt.figure()
plt.plot(df['Study_Time_Hours'], df['Exam_Score'], 'o')
plt.xlabel("Study Time (Hours)")
plt.ylabel("Exam Score")
plt.title("Line Plot: Study Time vs Exam Score")
plt.show()

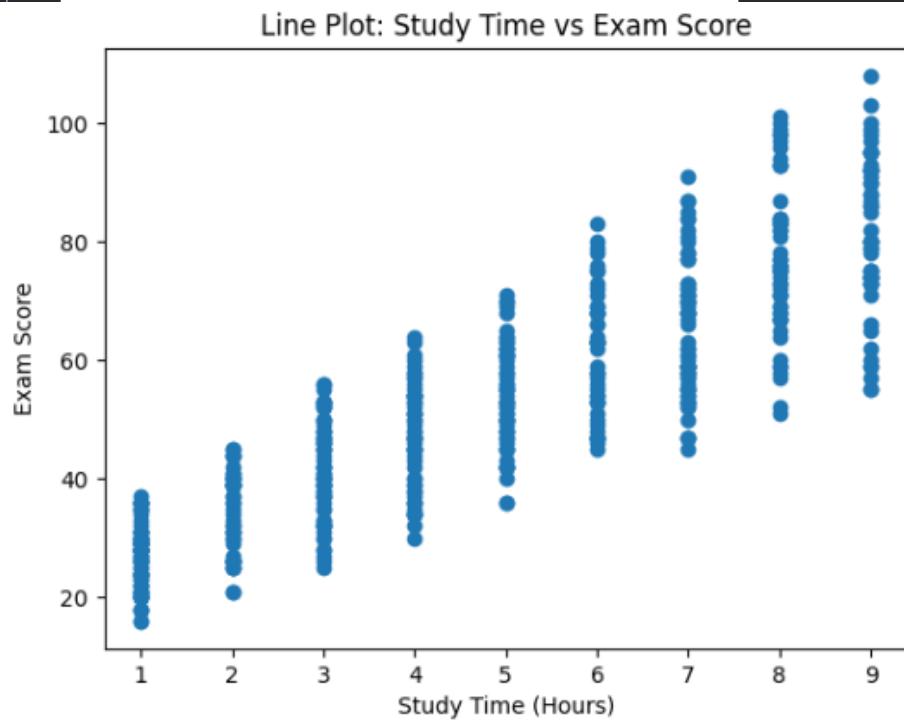
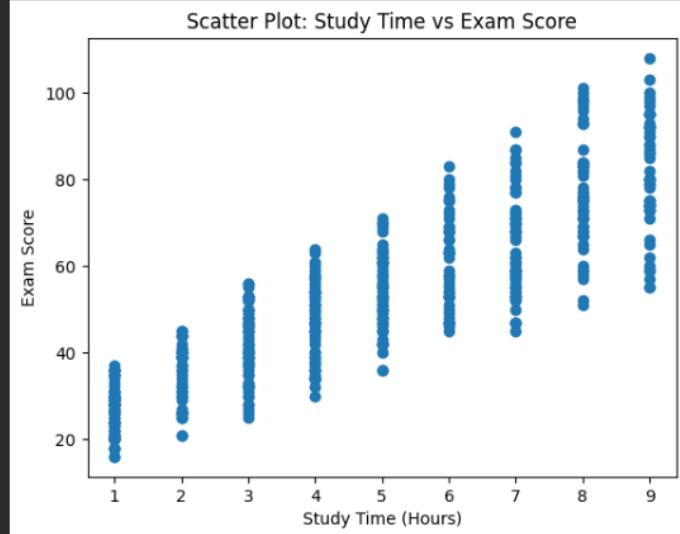
plt.figure()
df.boxplot(column='Exam_Score', by='Study_Time_Hours')
plt.xlabel("Study Time (Hours)")
plt.ylabel("Exam Score")
plt.title("Box Plot: Exam Score Distribution by Study Time")
plt.suptitle("")
plt.show()

if correlation > 0.5:
    print("\nInsight: Strong positive correlation - more study time generally leads to higher scores.")
elif correlation > 0:
    print("\nInsight: Moderate positive correlation between study time and exam score.")
else:
    print("\nInsight: Weak or no correlation between study time and exam score.")

```

```
...   Study_Time_Hours  Exam_Score  
0           7          85  
1           4          45  
2           8          98  
3           5          56  
4           7          59
```

```
Correlation between Study Time and Exam Score: 0.8724061772649271
```



EXP_38

```
▶ import pandas as pd
df = pd.read_csv("city_temperature_data.csv")

print(df.head())

mean_temp = df.groupby('City')[ 'Temperature'].mean()
print("\nMean Temperature for Each City:")
print(mean_temp)

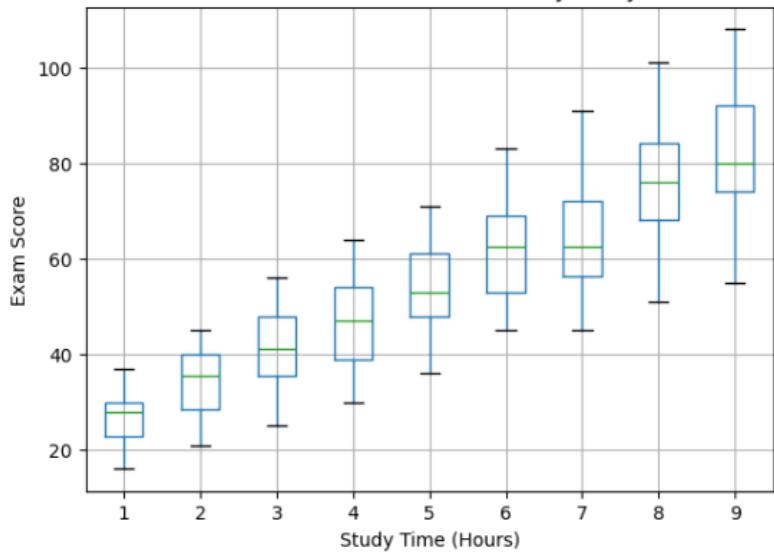
std_temp = df.groupby('City')[ 'Temperature'].std()
print("\nStandard Deviation of Temperature for Each City:")
print(std_temp)

temp_range = df.groupby('City')[ 'Temperature'].max() - df.groupby('City')[ 'Temperature'].min()
print("\nTemperature Range for Each City:")
print(temp_range)
highest_range_city = temp_range.idxmax()
print("\nCity with Highest Temperature Range:", highest_range_city)

most_consistent_city = std_temp.idxmin()
print("City with Most Consistent Temperature:", most_consistent_city)
```

<Figure size 640x480 with 0 Axes>

Box Plot: Exam Score Distribution by Study Time



Insight: Strong positive correlation – more study time generally leads to higher scores.

EXP_39

```
▶ import pandas as pd
df = pd.read_csv("city_temperature_data.csv")

print(df.head())

mean_temp = df.groupby('City')[ 'Temperature'].mean()
print("\nMean Temperature for Each City:")
print(mean_temp)

std_temp = df.groupby('City')[ 'Temperature'].std()
print("\nStandard Deviation of Temperature for Each City:")
print(std_temp)

temp_range = df.groupby('City')[ 'Temperature'].max() - df.groupby('City')[ 'Temperature'].min()
print("\nTemperature Range for Each City:")
print(temp_range)
highest_range_city = temp_range.idxmax()
print("\nCity with Highest Temperature Range:", highest_range_city)

most_consistent_city = std_temp.idxmin()
print("City with Most Consistent Temperature:", most_consistent_city)
```

```
...
   City Day  Temperature
0  Chennai    1      32.48
1  Chennai    2      29.31
2  Chennai    3      33.24
3  Chennai    4      37.62
4  Chennai    5      28.83
```

Mean Temperature for Each City:

```
City
Bengaluru    30.178000
Chennai      30.049890
Delhi        29.810219
Kolkata      30.408548
Mumbai        30.634603
Name: Temperature, dtype: float64
```

Standard Deviation of Temperature for Each City:

```
City
Bengaluru    4.992915
Chennai      4.740282
Delhi        5.097348
Kolkata      5.106776
Mumbai        4.784817
Name: Temperature, dtype: float64
```

Temperature Range for Each City:

```
City
Bengaluru    30.58
Chennai      35.47
Delhi        28.87
Kolkata      27.49
Mumbai        27.40
Name: Temperature, dtype: float64
```

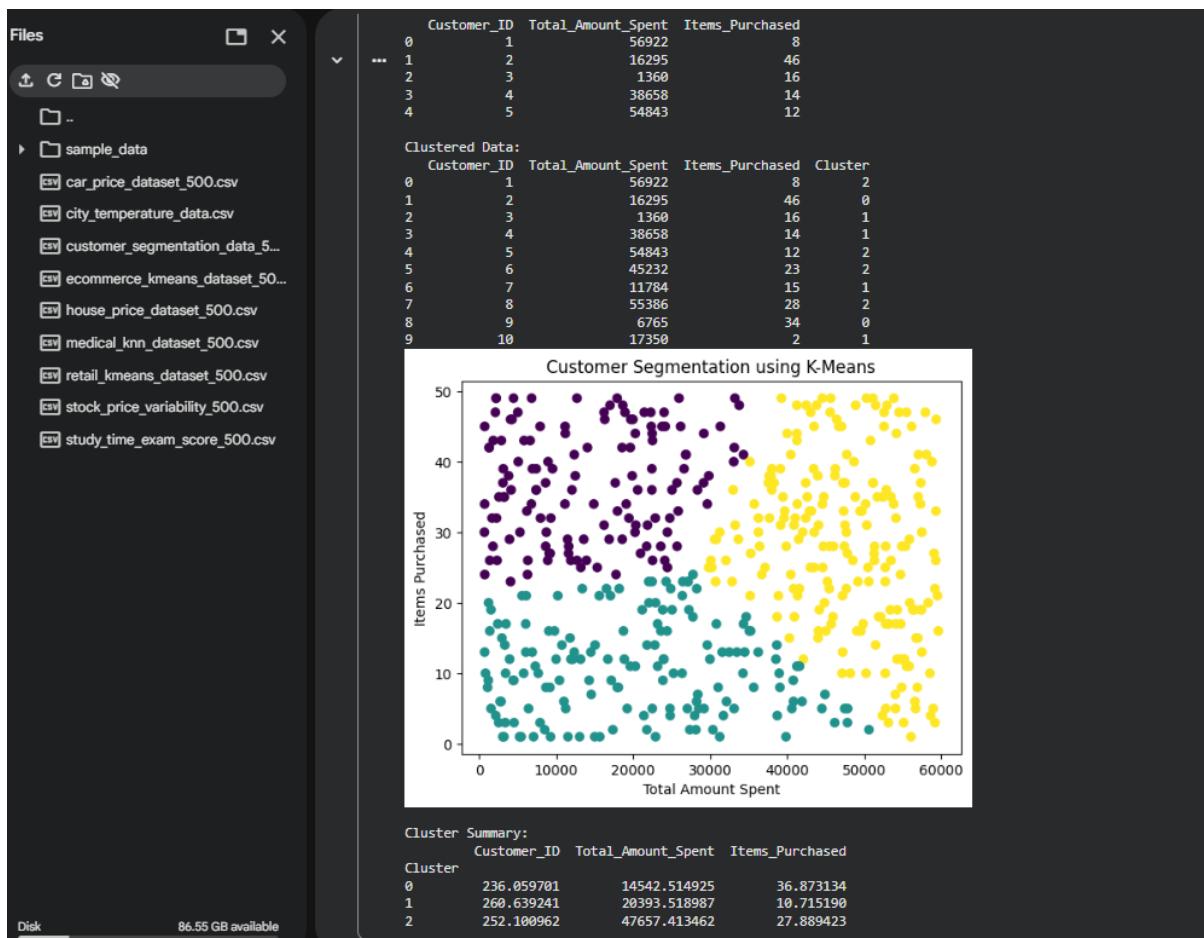
```
City with Highest Temperature Range: Chennai
City with Most Consistent Temperature: Chennai
```

City with Most Consistent Temperature: Chennai

```
[10] ✓ 1s
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

df = pd.read_csv("ecommerce_kmeans_dataset_500.csv")

print(df.head())
X = df[['Total_Amount_Spent', 'Items_Purchased']]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
kmeans = KMeans(n_clusters=3, random_state=42)
df['Cluster'] = kmeans.fit_predict(X_scaled)
print("\nClustered Data:")
print(df.head(10))
plt.scatter(
    df['Total_Amount_Spent'],
    df['Items_Purchased'],
    c=df['Cluster']
)
plt.xlabel("Total Amount Spent")
plt.ylabel("Items Purchased")
plt.title("Customer Segmentation using K-Means")
plt.show()
cluster_summary = df.groupby('Cluster').mean()
print("\nCluster Summary:")
print(cluster_summary)
```



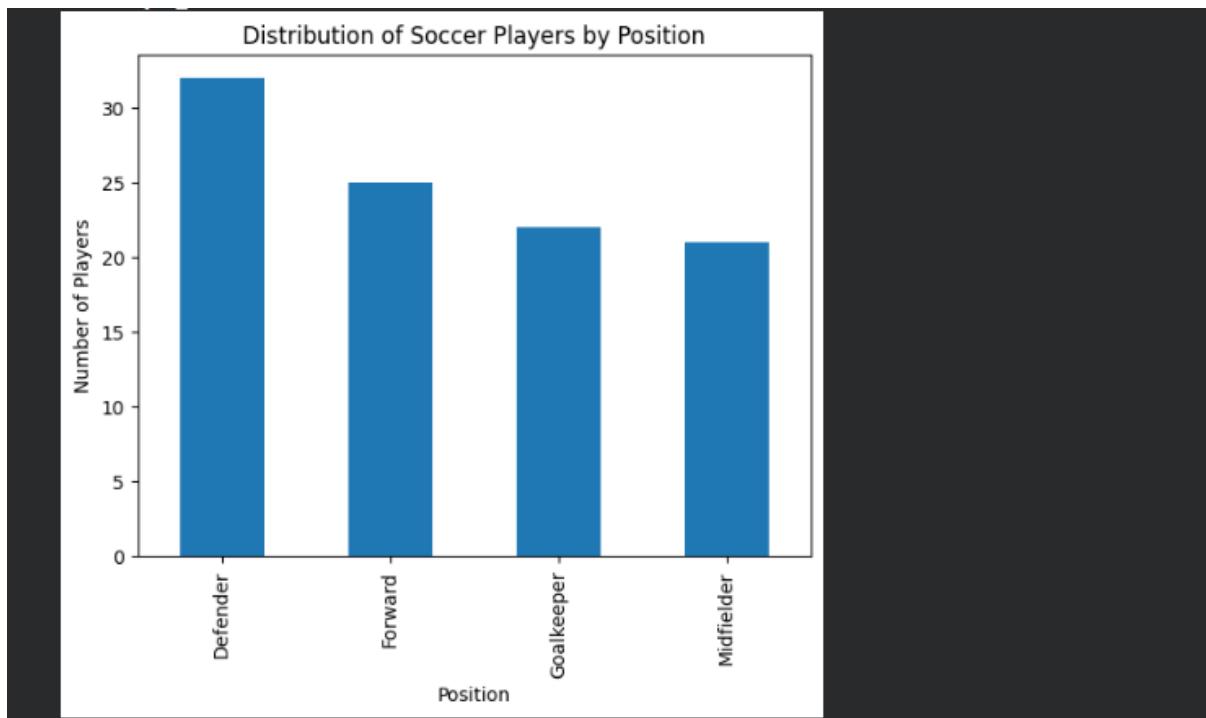
EXP_40

```
[1]: house_price_dataset_500.csv  
[2]: medical_knn_dataset_500.csv  
[3]: retail_kmeans_dataset_500.csv  
[4]: soccer_players_data.csv  
[5]: stock_price_variability_500.csv  
[6]: study_time_exam_score_500.csv
```

```
[1]:  
import pandas as pd  
import matplotlib.pyplot as plt  
df = pd.read_csv("soccer_players_data.csv")  
  
print(df.head())  
top_goals = df.sort_values(by="Goals_Scored", ascending=False).head(5)  
print("\nTop 5 Players with Highest Goals:")  
print(top_goals[['Player_Name', 'Goals_Scored']])  
top_salary = df.sort_values(by="Weekly_Salary", ascending=False).head(5)  
print("\nTop 5 Players with Highest Salaries:")  
print(top_salary[['Player_Name', 'Weekly_Salary']])  
average_age = df['Age'].mean()  
print("\nAverage Age of Players:", average_age)  
above_avg_age = df[df['Age'] > average_age]  
print("\nPlayers Above Average Age:")  
print(above_avg_age[['Player_Name', 'Age']])  
position_count = df['Position'].value_counts()  
  
plt.figure()  
position_count.plot(kind='bar')  
plt.xlabel("Position")  
plt.ylabel("Number of Players")  
plt.title("Distribution of Soccer Players by Position")  
plt.show()
```

```
[...]  
...  
    Player_Name  Age   Position  Goals_Scored  Weekly_Salary  
0   Player_1    24   Defender      31        342890  
1   Player_2    37   Goalkeeper     23        99377  
2   Player_3    32   Goalkeeper     11        491620  
3   Player_4    28   Midfielder     1         366189  
4   Player_5    25   Defender       2        410032  
  
Top 5 Players with Highest Goals:  
Player_Name  Goals_Scored  
75   Player_76      34  
73   Player_74      34  
52   Player_53      34  
62   Player_63      34  
81   Player_82      34  
  
Top 5 Players with Highest Salaries:  
Player_Name  Weekly_Salary  
88   Player_89      498982  
33   Player_34      497456  
48   Player_49      495101  
72   Player_73      494209  
70   Player_71      492905  
  
Average Age of Players: 27.01
```

```
... Players Above Average Age:  
Player_Name  Age  
1   Player_2   37  
2   Player_3   32  
3   Player_4   28  
6   Player_7   36  
7   Player_8   28  
8   Player_9   28  
13  Player_14  29  
17  Player_18  29  
18  Player_19  29  
19  Player_20  34  
21  Player_22  33  
22  Player_23  32  
23  Player_24  32  
24  Player_25  36  
25  Player_26  29  
26  Player_27  37  
29  Player_30  36  
33  Player_34  35  
35  Player_36  31  
36  Player_37  35  
39  Player_40  37  
40  Player_41  32  
42  Player_43  29  
44  Player_45  32  
46  Player_47  31  
47  Player_48  34  
49  Player_50  35  
56  Player_57  35  
57  Player_58  29  
61  Player_62  31  
62  Player_63  33  
63  Player_64  32  
65  Player_66  31  
67  Player_68  33  
68  Player_69  30  
69  Player_70  35  
70  Player_71  32  
71  Player_72  30  
73  Player_74  32  
74  Player_75  30  
79  Player_80  29  
81  Player_82  28  
82  Player_83  36  
83  Player_84  34  
95  Player_96  29  
98  Player_99  33
```



Settings

Languages & Frameworks > Android SDK

Manager for the Android SDK and Tools used by the IDE

Android SDK Location: [Edit](#) Optimize disk space

SDK Platforms **SDK Tools** **SDK Update Sites**

The Android SDK location cannot be at the filesystem root.

Each Android SDK Platform package includes the Android platform and sources pertaining to an API level by default. Once installed, the IDE will automatically check for updates. Check "show package details" to display individual SDK components.

Name	API Level	Revision	Status
<input type="checkbox"/> Android 14.0 ("UpsideDownCake")	34-ext10	1	Not installed
<input type="checkbox"/> Android 14.0 ("UpsideDownCake")	34-ext11	1	Not installed
<input type="checkbox"/> Android 14.0 ("UpsideDownCake")	34-ext12	1	Not installed
<input type="checkbox"/> Android UpsideDownCake Preview	UpsideDownCake	4	Not installed
<input type="checkbox"/> Android 13.0 ("Tiramisu")	33	3	Not installed
<input type="checkbox"/> Android 13.0 ("Tiramisu")	33-ext4	1	Not installed
<input type="checkbox"/> Android 13.0 ("Tiramisu")	33-ext5	1	Not installed
<input type="checkbox"/> Android 12L ("Sv2")	32	1	Not installed
<input type="checkbox"/> Android 12.0 ("S")	31	1	Not installed
<input type="checkbox"/> Android 11.0 ("R")	30	3	Not installed
<input type="checkbox"/> Android 10.0 ("Q")	29	5	Not installed
<input type="checkbox"/> Android 9.0 ("Pie")	28	6	Not installed
<input type="checkbox"/> Android 8.1 ("Oreo")	27	3	Not installed

Hide Obsolete Packages Show Package Details

Project-level settings will be applied to new projects **OK** **Cancel** **Apply**